

Towards Exascale with IPPL Version 2.0

OPAL developers meeting

Matthias Frey¹ Sriramkrishnan Muralikrishnan² Andreas Adelmann²

¹Mathematical Institute, University of St Andrews, UK

²Paul Scherrer Institute, PSI Villigen, CH

19 Jan 2022

IPPL 2.0 developers



- ▶ **Portability** across serial, distributed, and parallel architectures
- ▶ Development of reusable, cross-problem-domain components to **enable rapid application development**
- ▶ **High efficiency** for kernels and components relevant to scientific simulation
- ▶ Framework **design and development driven by applications** from a diverse set of scientific problem domains
- ▶ **Shorter time** from problem inception **to working parallel simulations**

Performance Portable Abstraction Layers

library	lang	no. contr. ³	license	repository
Kokkos	C++	99	BSD 3-Clause	https://github.com/kokkos
alpaka	C++	28	MPL-2.0	https://github.com/alpaka-group/alpaka
Thrust	C++	65	Apache License 2.0	https://github.com/thrust/thrust
RAJA	C++	42	BSD 3-Clause	https://github.com/LLNL/RAJA

backend	Kokkos	alpaka	Thrust	RAJA
OpenMP	✓	✓	✓	✓
Pthreads	✓	✗	✗	✗
C++11 threads	✗	✓	✗	✗
CUDA	✓	✓	✓	✓
TBB	✗	✓	✓	⚠️
HPX	✓	✗	✗	✗
HIP ⁴	✓	✓	✗	✓
SYCL ⁵	✓	✗	✗	⚠️

programming	Kokkos	alpaka	Thrust	RAJA
memory management complexity	✓ medium/high	✓ high	✓ low/medium	✗ medium/high

³Number of contributors on 18 Jan 2022

⁴Interface for NVIDIA and AMD GPUs

⁵Interface for NVIDIA, AMD and Intel CPUs and GPUs

- ▶ Enable **performance portability with Kokkos**
(i.e. replace field and particle containers with Kokkos data structures)
- ▶ Upgrade to **C++17 standard**
- ▶ Keep **expression templates** for particles and fields
- ▶ **Keep** changes to the **user interface** to a minimum
- ▶ **Simplify code**



► Attributes:

- ▶ Struct of Kokkos::Views
- ▶ Expression templates
- ▶ Easily added to application

► Communication:

- ▶ Particle layout classes
- ▶ De-/serialize Kokkos::View<char*>
- ▶ Pre-allocated buffers

```
using namespace ippl;


template<class PLayout>
struct Bunch
: public ParticleBase<PLayout> {
    Bunch(PLayout& playout)
    : ParticleBase<PLayout>(playout)
    {
        // add application attributes
        this->addAttribute(R);
        this->addAttribute(V);
        this->addAttribute(mass);
        this->addAttribute(charge);
    }

    ~Bunch() { }

    ParticleAttrib<double> mass, charge;

    ParticleAttrib<Vector<double>> R, V;
};

// compiles to single Kokkos kernel
bunch->R = bunch->R + dt * bunch->V;
```



▶ Scalar/Vector fields:

- ▶ Kokkos::Views
- ▶ Expression templates

▶ Interface to heFFTe⁶

▶ Communication:

- ▶ Field layout
- ▶ Distribution of local domains globally known
- ▶ De-/serialize Kokkos::View<char*>
- ▶ Pre-allocated buffers

```
using namespace ippl;

constexpr unsigned int dim = 3;

int pt = 256;
Index I(pt);
NDIndex<dim> owned(I, I, I);

// specifies SERIAL, PARALLEL dimensions
e_dim_tag decomp[dim] = {PARALLEL,
                           PARALLEL,
                           PARALLEL};

FieldLayout<dim> layout(owned, decomp);

double dx = 1.0 / double(pt);
Vector<double, dim> hx = {dx, dx, dx};
Vector<double, dim> origin = {0, 0, 0};

UniformCartesian<double, dim>
    mesh(owned, hx, origin);

Field<double, dim> field(mesh, layout);
```

⁶Ayala A., Tomov S., Haidar A., Dongarra J. (2020) heFFTe: Highly Efficient FFT for Exascale. ICCS 2020.
https://doi.org/10.1007/978-3-030-50371-0_19


Expression templates

- ▶ Avoids temporary objects in mathematical expressions
- ▶ Reduces number of Kokkos kernels
- ▶ Assignment operator evaluates expression
- ▶ Available for particles and fields
- ▶ Supported:
 - ▶ Binary operations: $+, -, *, /$
 - ▶ Comparison operations: $<, \leq, >, \geq, ==, !=$
 - ▶ Bitwise operations
 - ▶ Many operations of cmath header

```
template<typename T, class... Properties>
template <typename E, size_t N>
ParticleAttrib<T, Properties...>&
ParticleAttrib<T, Properties...>::operator=(  
    detail::Expression<E, N> const& expr)  
{  
    using capture_type =  
        detail::CapturedExpression<E, N>;  
  
    capture_type expr_ =  
        reinterpret_cast<  
            const capture_type&>(expr);  
  
    Kokkos::parallel_for(  
        "ParticleAttrib::operator()",  
        dview_m.extent(0),  
        KOKKOS_CLASS_LAMBDA(const size_t i) {  
            dview_m(i) = expr_(i);  
        });  
  
    return *this;  
}
```


Particles:

- ▶ Locate particles to send
- ▶ Pack attributes into buffer
- ▶ Send buffer to receiver
- ▶ Unpack attributes



Fields:




- ▶ Get grid intersection
- ▶ Pack intersection into buffer
- ▶ Send buffer to receiver
- ▶ Unpack intersection



Domain Decomposition – Orthogonal Recursive Bisection

ETH Semester project, Michael Ligotino


1. Interpolate particles to grid (weights only)
2. Slice cut axis and evaluate sum of weights S_i
3. Evaluate median over all S_i
4. Cut (sub-)domain at median



Example with 64 MPI ranks

Conjugate Gradient Solver – Timings for 512^3 grid

ETH Bachelor thesis, Alessandro Vinciguerra





FFT Poisson Solver for solving Hose instability

ETH Master thesis, Sonali Mayani

Algorithms:

- ▶ Hockney-Eastwood⁷
- ▶ Vico-Greengard⁸





⁷Hockney, R. and Eastwood, J. (1988). Computer Simulation Using Particles. CRC Press.

⁸Vico, F., Greengard, L., and Ferrando, M. (2016). Fast convolution with free-space green's functions. Journal of Computational Physics, 323:191–203.


ALPINE: A set of portable pLasma physics Particle-in-cell mINI-apps for Exascale

Two-stream Instability Penning trap Landau damping

- Light weight codes
- Proxy for real applications
- For implementing new algorithms
- Testing new implementations
- Reference results ensure correctness




OPAL: Object oriented Parallel Accelerator Library
Open source C++ Library for particle accelerators being developed by twelve core developers across seven institutes




Source: OPAL web page

Independent Parallel Particle Layer (IPPL) v 2.0




Linear Landau damping: Strong scaling

time/time step (s)




Total time



Efficiency

Linear Landau damping: Weak scaling



- ▶ **For GPUs:** 256×128^2 grid and 8 particles per cell is the base case for 1 node/GPU. The max. grid size and particles are 2048^3 and $Np \approx 69$ billion at 2048 GPUs
- ▶ **For CPUs:** 512×256^2 grid and 8 particles per cell is the base case for 1 node. The max. grid size and particles are 4096×2048^2 and $Np \approx 138$ billion at 512 nodes = 16,384 cores

- ▶ Improve load balancing and memory usage
- ▶ Use IPPL 2.0 in OPAL ("partial rewrite OPAL required")

Thank you for your attention!

Thanks to

- ▶ Alessandro Vinciguerra
- ▶ Sonali Mayani
- ▶ Michael Ligotino